

- Which of the following best explains how multi-head attention improves model understanding in Transformers?
 - By reducing the total number of parameters through parallelization
 - By enforcing uniform attention over the sequence to prevent bias
 - By increasing computation speed through batch-wise attention
 - By enabling different heads to attend to diverse relational patterns across positions
- Which component of the Transformer architecture is exclusively utilized in GPT, making it more suited for generative tasks?
 - Decoder layers with masked self-attention
 - Encoder layers for input sequence modeling
 - A hybrid encoder-decoder combination
 - A purely feed-forward architecture
- What design choice in GPT restricts it from leveraging full bidirectional context, and what consequence does this have?
 - Encoder-based design; restricts output generation
 - Unidirectional left-to-right flow; limits contextual understanding
 - Bidirectional masking; leads to complex fitting
 - Cross-attention dependencies; increase inference latency
- Which of the following best characterizes the training objectives that enable BERT to capture both deep token-level context and inter-sentence semantics?
 - Predicting the next token in a left-to-right fashion using unidirectional context
 - Learning to generate a target sequence from an input sequence in an encoder-decoder setup
 - Jointly optimizing masked token reconstruction and inter-sentence coherence discrimination
 - Aligning image features with textual descriptions through cross-modal supervision

Short Answer Questions

- What are the potential drawbacks of the two-stage process of pretraining on large corpora followed by fine-tuning on specific tasks in Transformer models?
- What are the potential drawbacks of GPT's autoregressive training objective when applied to tasks requiring holistic understanding of the text?
- BERT utilizes a masked language model (MLM) during pretraining. What is the primary challenge associated with the MLM approach, and how does it affect the model's downstream performance?
- GPT models are known for their unidirectional (left-to-right) processing. How does this design choice impact their performance on tasks like text generation compared to tasks like text classification?