

Reinforcement Learning — Boxed Questions (Large Font)

Q1. (Multi-select) Which algorithms belong to the policy-gradient family?

- (A) Proximal Policy Optimization (PPO)
- (B) Q-learning
- (C) REINFORCE
- (D) Deep Q-Network (DQN)

Q2. An environment is considered partially observable when:

- (A) The agent can observe the full state
- (B) Observations lack some hidden variables
- (C) The reward signal is stochastic
- (D) Transition dynamics are deterministic

Q3. In Q-learning, the key difference from SARSA is:

- (A) On-policy vs. off-policy nature
- (B) Use of neural networks
- (C) Continuous action space support
- (D) Model-based planning

Q4. The Advantage in Actor–Critic methods is calculated as:

- (A) State-value minus action-value
- (B) Action-value minus state-value
- (C) Reward minus entropy
- (D) Policy-gradient estimate

Q5. (Multi-select) Techniques to stabilize deep-RL training include:

- (A) Experience replay
- (B) Target networks
- (C) Layer normalization
- (D) Early stopping

Q6. A discount factor γ close to 0 emphasizes:

- (A) Immediate rewards
- (B) Long-term rewards
- (C) Exploration
- (D) Deterministic policies

Q7. The exploration–exploitation dilemma refers to:

(A) Balancing policy and value networks

(B) Choosing between trying new actions and using known good actions

(C) Data augmentation

(D) Hyper-parameter tuning