

RNNs and LSTMs

Multiple Choice Questions

- Q1.** What is the primary benefit of stacking multiple RNN layers (i.e., stacked RNNs)?
- A. Faster training
 - B. Lower memory usage
 - C. Better learning of hierarchical features
 - D. Simpler architecture
- Q2.** Which of the following is the main reason RNNs struggle with long-term dependencies?
- A. Overfitting
 - B. Vanishing gradients
 - C. Lack of non-linearity
 - D. Insufficient data
- Q3.** What differentiates an LSTM cell from a standard RNN cell?
- A. It uses ReLU instead of tanh
 - B. It introduces gates to control the flow of information
 - C. It has fewer parameters
 - D. It is a convolutional architecture
- Q4.** In a standard LSTM, which gate is responsible for deciding how much of the past memory to keep?
- A. Output gate
 - B. Forget gate
 - C. Input gate
 - D. Update gate

Descriptive Questions

- Q5.** Why is the forget gate bias in LSTMs often initialized to a high value (e.g., 2 or 3)? Explain its effect on long-term dependency learning.
- Q6.** Bidirectional RNNs are often used for POS tagging but not machine translation. Explain why, considering input-output alignment and context flow.
- Q7.** Designing an RNN model for variable-length legal documents with long dependencies:
- (a) Choose between vanilla RNN or LSTM.
 - (b) Stack layers or keep it shallow?
 - (c) Make it bidirectional?
- Justify each choice based on model behavior and task needs.
- Q8.** Consider a vanilla RNN with recurrent weight matrix W_h and sequence length 50. Analyze gradient behavior:
- (1) If $\|W_h\| = 0.9$: Will gradients vanish or explode? Justify.
 - (2) If $\|W_h\| = 1.2$: Will gradients vanish or explode? Justify. Suggest an easy fix and explain how it helps.
- Hint: Consider eigenvalue effects on gradient propagation over time.*